



МИНОБРНАУКИ РОССИИ

Федеральное государственное бюджетное образовательное учреждение
высшего образования

**«МИРЭА – Российский технологический университет»
РТУ МИРЭА**

ИКБ направление «Киберразведка и противодействие угрозам с применением технологий
искусственного интеллекта» 10.04.01

Кафедра КБ-4 «Интеллектуальные системы информационной безопасности»

Отчет по Практической работе №6 и Лабораторной работе №4

по дисциплине: «Анализ защищенности систем искусственного
интеллекта»

На тему: «Изучение методов защиты от атак на модели НС»

Группа:

ББМО-01-22

Выполнил:

Гребенник Г.С

Проверил:

к.т.н. Спирин А.А.

Москва, 202

Цель работы

Провести серию атак на НС, после чего реализовать функцию защиты и сделать вывод по полученному результату, оценив стойкость предложенного способа защиты.

Ход работы:

1. Импортируем необходимые библиотеки:

```
import numpy as np
import matplotlib.pyplot as plt
import torch
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
from torchvision import transforms, datasets
```

2. Задаем нормализующие преобразования, загружаем набор данных и разбиваем данные на подвыборки:

```
[2] transform = transforms.Compose([transforms.ToTensor(),
transforms.Normalize((0.0,), (1.0,))])
dataset = datasets.MNIST(root = './data', train=True, transform = transform, download=True)
train_set, val_set = torch.utils.data.random_split(dataset, [50000, 10000])
test_set = datasets.MNIST(root = './data', train=False, transform = transform, download=True)
train_loader = torch.utils.data.DataLoader(train_set, batch_size=1, shuffle=True)
val_loader = torch.utils.data.DataLoader(val_set, batch_size=1, shuffle=True)
test_loader = torch.utils.data.DataLoader(test_set, batch_size=1, shuffle=True)
print("Training data:", len(train_loader), "Validation data:", len(val_loader), "Test data:", len(test_loader))

Downloading http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz
to ./data/MNIST/raw/train-images-idx3-ubyte.gz
100%|██████████| 9912422/9912422 [00:00<00:00, 82111399.50it/s]
Extracting ./data/MNIST/raw/train-images-idx3-ubyte.gz to ./data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz
to ./data/MNIST/raw/train-labels-idx1-ubyte.gz
100%|██████████| 28881/28881 [00:00<00:00, 89465061.91it/s]
Extracting ./data/MNIST/raw/train-labels-idx1-ubyte.gz to ./data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz
to ./data/MNIST/raw/t10k-images-idx3-ubyte.gz
100%|██████████| 1648877/1648877 [00:00<00:00, 31326511.98it/s]
Extracting ./data/MNIST/raw/t10k-images-idx3-ubyte.gz to ./data/MNIST/raw

Downloading http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz
to ./data/MNIST/raw/t10k-labels-idx1-ubyte.gz
100%|██████████| 4542/4542 [00:00<00:00, 21821911.53it/s]Extracting ./data/MNIST/raw/t10k-labels-idx1-ubyte.gz to ./data/MNIST/raw

Training data: 50000 Validation data: 10000 Test data: 10000
```

3. Настраиваем использование графического ускорителя:

```
use_cuda=True
device = torch.device("cuda" if (use_cuda and torch.cuda.is_available()) else "cpu")
```

Создание атак на модель НС

1. Создаем класс НС на основе torch:

```

✓
0
CEK.
[4] class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(1, 32, 3, 1)
        self.conv2 = nn.Conv2d(32, 64, 3, 1)
        self.dropout1 = nn.Dropout2d(0.25)
        self.dropout2 = nn.Dropout2d(0.5)
        self.fc1 = nn.Linear(9216, 128)
        self.fc2 = nn.Linear(128, 10)
    def forward(self, x):
        x = self.conv1(x)
        x = F.relu(x)
        x = self.conv2(x)
        x = F.relu(x)
        x = F.max_pool2d(x, 2)
        x = self.dropout1(x)
        x = torch.flatten(x, 1)
        x = self.fc1(x)
        x = F.relu(x)
        x = self.dropout2(x)
        x = self.fc2(x)
        output = F.log_softmax(x, dim=1)
        return output

```

2. Проверка работоспособности созданного класса НС:

```

✓
0
CEK.
[5] model = Net().to(device)

```

3. Создадим оптимизатор, функцию потерь и трейнер сети:

```

✓
0
CEK.
▶ optimizer = optim.Adam(model.parameters(), lr=0.0001, betas=(0.9, 0.999))
  criterion = nn.NLLLoss()
  scheduler = optim.lr_scheduler.ReduceLROnPlateau(optimizer, mode='min', factor=0.1, patience=3)

```

4. Определим функцию обучения сети:

```

[7] def fit(model,device,train_loader,val_loader,epochs):
    data_loader = {'train':train_loader,'val':val_loader}
    print("Fitting the model...")
    train_loss,val_loss=[],[]
    for epoch in range(epochs):
        loss_per_epoch,val_loss_per_epoch=0,0
        for phase in ('train','val'):
            for i,data in enumerate(data_loader[phase]):
                input,label = data[0].to(device),data[1].to(device)
                output = model(input)
                #calculating loss on the output
                loss = criterion(output,label)
                if phase == 'train':
                    optimizer.zero_grad()
                    #grad calc w.r.t Loss func
                    loss.backward()
                    #update weights
                    optimizer.step()
                    loss_per_epoch+=loss.item()
                else:
                    val_loss_per_epoch+=loss.item()
            scheduler.step(val_loss_per_epoch/len(val_loader))
        print("Epoch: {} Loss: {} Val_Loss: {}".format(epoch+1,loss_per_epoch/len(train_loader),val_loss_per_epoch/len(val_loader)))
        train_loss.append(loss_per_epoch/len(train_loader))
        val_loss.append(val_loss_per_epoch/len(val_loader))
    return train_loss,val_loss

```

5. Обучим модель:

```

loss, val_loss = fit(model, device, train_loader, val_loader, 10)


```

Fitting the model...

/usr/local/lib/python3.10/dist-packages/torch/nn/functional.py:1345: UserWarning: dropout2d: Received a 2-D input to dropout2d, which is deprecated and will result in an error in a future version of PyTorch. Please use dropout instead. (Triggered at: /usr/local/lib/python3.10/dist-packages/torch/nn/modules/module.py:1154: Module.__getattr__: 'dropout' is deprecated and will be removed in a future version of PyTorch.)

Epoch: 1 Loss: 0.2850503676612788 Val_Loss: 0.1465472458492094
Epoch: 2 Loss: 0.11141908125931078 Val_Loss: 0.11816810167795906
Epoch: 3 Loss: 0.08508949297238538 Val_Loss: 0.10610676416118986
Epoch: 4 Loss: 0.07273402502575843 Val_Loss: 0.08871200382142519
Epoch: 5 Loss: 0.0640585860067484 Val_Loss: 0.08652106153088714
Epoch: 6 Loss: 0.0607716777421671 Val_Loss: 0.08539433175660129
Epoch: 7 Loss: 0.0562644432195245 Val_Loss: 0.07379816616976491
Epoch: 8 Loss: 0.053677760184325574 Val_Loss: 0.07861828022329812
Epoch: 9 Loss: 0.052243182574612974 Val_Loss: 0.07426569758173895
Epoch: 10 Loss: 0.050123429469401834 Val_Loss: 0.07791693960269636

6. Построим графики потерь при обучении и валидации в зависимости от эпохи:



7. Создаем функции атак FGSM, I-FGSM, MI-FGSM:

✓
0
CEK.

```
def fgsm_attack(input,epsilon,data_grad):  
    pert_out = input + epsilon*data_grad.sign()  
    pert_out = torch.clamp(pert_out, 0, 1)  
    return pert_out
```

✓
0
CEK.

```
[11] def ifgsm_attack(input,epsilon,data_grad):  
    iter = 10  
    alpha = epsilon/iter  
    pert_out = input  
    for i in range(iter-1):  
        pert_out = pert_out + alpha*data_grad.sign()  
        pert_out = torch.clamp(pert_out, 0, 1)  
        if torch.norm((pert_out-input),p=float('inf')) > epsilon:  
            break  
    return pert_out
```

✓
0
CEK.

```
[12] def mifgsm_attack(input,epsilon,data_grad):  
    iter=10  
    decay_factor=1.0  
    pert_out = input  
    alpha = epsilon/iter  
    g=0  
    for i in range(iter-1):  
        g = decay_factor*g + data_grad/torch.norm(data_grad,p=1)  
        pert_out = pert_out + alpha*torch.sign(g)  
        pert_out = torch.clamp(pert_out, 0, 1)  
        if torch.norm((pert_out-input),p=float('inf')) > epsilon:  
            break  
    return pert_out
```

8. Создаем функцию проверки:

```

0
CEK.
def test(model, device, test_loader, epsilon, attack):
    correct = 0
    adv_examples = []
    for data, target in test_loader:
        data, target = data.to(device), target.to(device)
        data.requires_grad = True
        output = model(data)
        init_pred = output.max(1, keepdim=True)[1]
        if init_pred.item() != target.item():
            continue
        loss = F.nll_loss(output, target)
        model.zero_grad()
        loss.backward()
        data_grad = data.grad.data
        if attack == "fgsm":
            perturbed_data = fgsm_attack(data, epsilon, data_grad)
        elif attack == "ifgsm":
            perturbed_data = ifgsm_attack(data, epsilon, data_grad)
        elif attack == "mifgsm":
            perturbed_data = mifgsm_attack(data, epsilon, data_grad)
        output = model(perturbed_data)
        final_pred = output.max(1, keepdim=True)[1]
        if final_pred.item() == target.item():
            correct += 1
        if (epsilon == 0) and (len(adv_examples) < 5):
            adv_ex = perturbed_data.squeeze().detach().cpu().numpy()
            adv_examples.append( (init_pred.item(), final_pred.item(), adv_ex) )
        else:
            if len(adv_examples) < 5:
                adv_ex = perturbed_data.squeeze().detach().cpu().numpy()
                adv_examples.append( (init_pred.item(), final_pred.item(), adv_ex) )
    final_acc = correct/float(len(test_loader))
    print("Epsilon: {} \t Test Accuracy = {} / {} = {}".format(epsilon, correct, len(test_loader), final_acc))
    return final_acc, adv_examples

```


9. Построим графики успешности атак и примеры выполненных атак в зависимости от степени возмущения epsilon:



































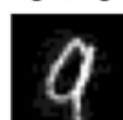





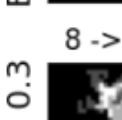
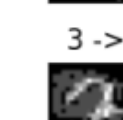
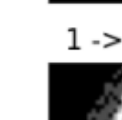
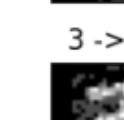
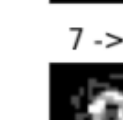
```

▶ epsilons = [0,0.007,0.01,0.02,0.03,0.05,0.1,0.2,0.3]
for attack in ("fgsm","ifgsm","mifgsm"):
    accuracies = []
    examples = []
    for eps in epsilons:
        acc, ex = test(model, device, test_loader, eps, attack)
        accuracies.append(acc)
        examples.append(ex)
    plt.figure(figsize=(5,5))
    plt.plot(epsilons, accuracies, "*-")
    plt.title(attack)
    plt.xlabel("Epsilon")
    plt.ylabel("Accuracy")
    plt.show()
    cnt = 0
    plt.figure(figsize=(8,10))
    for i in range(len(epsilons)):
        for j in range(len(examples[i])):
            cnt += 1
            plt.subplot(len(epsilons),len(examples[0]),cnt)
            plt.xticks([], [])
            plt.yticks([], [])
            if j == 0:
                plt.ylabel("Eps: {}".format(epsilons[i]), fontsize=14)
            orig,adv,ex = examples[i][j]
            plt.title("{} -> {}".format(orig, adv))
            plt.imshow(ex, cmap="gray")
    plt.tight_layout()
    plt.show()


```







































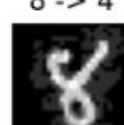






Epsilon: 0 Test Accuracy = 9730 / 10000 = 0.973
Epsilon: 0.007 Test Accuracy = 9706 / 10000 = 0.9706
Epsilon: 0.01 Test Accuracy = 9674 / 10000 = 0.9674
Epsilon: 0.02 Test Accuracy = 9650 / 10000 = 0.965
Epsilon: 0.03 Test Accuracy = 9592 / 10000 = 0.9592
Epsilon: 0.05 Test Accuracy = 9436 / 10000 = 0.9436
Epsilon: 0.1 Test Accuracy = 8778 / 10000 = 0.8778
Epsilon: 0.2 Test Accuracy = 6428 / 10000 = 0.6428
Epsilon: 0.3 Test Accuracy = 4143 / 10000 = 0.4143




	0 -> 0	9 -> 9	2 -> 2	3 -> 3	8 -> 8
Eps: 0					
	9 -> 9	6 -> 6	4 -> 4	0 -> 0	4 -> 4
Eps: 0.007					
	8 -> 8	3 -> 3	2 -> 2	7 -> 7	8 -> 8
Eps: 0.01					
	2 -> 2	4 -> 4	1 -> 1	3 -> 3	6 -> 6
Eps: 0.02					
	9 -> 9	2 -> 2	4 -> 4	5 -> 5	3 -> 3
Eps: 0.03					
	3 -> 3	6 -> 6	3 -> 3	4 -> 4	6 -> 6
Eps: 0.05					
	8 -> 8	1 -> 1	3 -> 3	7 -> 7	9 -> 9
Eps: 0.1					
	0 -> 0	5 -> 5	1 -> 1	3 -> 3	6 -> 0
Eps: 0.2					
	8 -> 3	3 -> 6	1 -> 8	3 -> 3	7 -> 7
Eps: 0.3					
























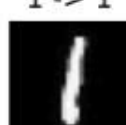







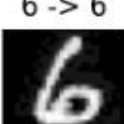
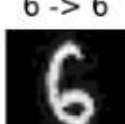
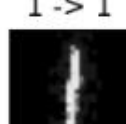






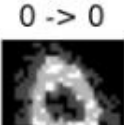
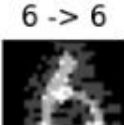

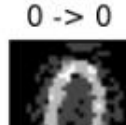

Epsilon: 0	Test Accuracy = 9729 / 10000 = 0.9729
Epsilon: 0.007	Test Accuracy = 9709 / 10000 = 0.9709
Epsilon: 0.01	Test Accuracy = 9699 / 10000 = 0.9699
Epsilon: 0.02	Test Accuracy = 9659 / 10000 = 0.9659
Epsilon: 0.03	Test Accuracy = 9606 / 10000 = 0.9606
Epsilon: 0.05	Test Accuracy = 9477 / 10000 = 0.9477
Epsilon: 0.1	Test Accuracy = 8936 / 10000 = 0.8936
Epsilon: 0.2	Test Accuracy = 6912 / 10000 = 0.6912
Epsilon: 0.3	Test Accuracy = 4839 / 10000 = 0.4839



Eps: 0	3 -> 3 	8 -> 8 	6 -> 6 	7 -> 7 	2 -> 2 
Eps: 0.007	0 -> 0 	9 -> 9 	7 -> 7 	1 -> 1 	3 -> 3 
Eps: 0.01	9 -> 9 	8 -> 8 	7 -> 7 	6 -> 6 	7 -> 7 
Eps: 0.02	2 -> 2 	4 -> 4 	2 -> 2 	3 -> 3 	2 -> 2 
Eps: 0.03	5 -> 5 	6 -> 6 	6 -> 6 	3 -> 3 	3 -> 3 
Eps: 0.05	6 -> 6 	3 -> 3 	8 -> 8 	0 -> 0 	6 -> 4 
Eps: 0.1	1 -> 1 	4 -> 4 	4 -> 4 	7 -> 2 	0 -> 0 
Eps: 0.2	3 -> 3 	4 -> 4 	7 -> 7 	8 -> 4 	8 -> 8 
Eps: 0.3	7 -> 7 	6 -> 6 	0 -> 0 	2 -> 2 	5 -> 3 

Epsilon: 0	Test Accuracy = 9708 / 10000 = 0.9708
Epsilon: 0.007	Test Accuracy = 9706 / 10000 = 0.9706
Epsilon: 0.01	Test Accuracy = 9694 / 10000 = 0.9694
Epsilon: 0.02	Test Accuracy = 9646 / 10000 = 0.9646
Epsilon: 0.03	Test Accuracy = 9601 / 10000 = 0.9601
Epsilon: 0.05	Test Accuracy = 9494 / 10000 = 0.9494
Epsilon: 0.1	Test Accuracy = 8953 / 10000 = 0.8953
Epsilon: 0.2	Test Accuracy = 6999 / 10000 = 0.6999
Epsilon: 0.3	Test Accuracy = 4797 / 10000 = 0.4797



	8 -> 8	4 -> 4	3 -> 3	5 -> 6	0 -> 0
Eps: 0					
Eps: 0.007	8 -> 8 	9 -> 9 	8 -> 8 	9 -> 9 	8 -> 8 
Eps: 0.01	7 -> 7 	0 -> 0 	2 -> 2 	8 -> 8 	3 -> 3 
Eps: 0.02	8 -> 8 	1 -> 1 	6 -> 6 	2 -> 2 	1 -> 1 
Eps: 0.03	3 -> 3 	5 -> 5 	8 -> 8 	1 -> 1 	3 -> 3 
Eps: 0.05	0 -> 0 	6 -> 6 	4 -> 4 	0 -> 9 	3 -> 3 
Eps: 0.1	3 -> 3 	6 -> 6 	6 -> 6 	1 -> 1 	3 -> 3 
Eps: 0.2	2 -> 3 	9 -> 4 	5 -> 8 	2 -> 8 	8 -> 8 
Eps: 0.3	0 -> 0 	6 -> 6 	0 -> 6 	0 -> 0 	0 -> 2 

Защита от атак

1. Создаем 2 класса НС

```

✓ 0 class NetF(nn.Module):
  def __init__(self):
    super(NetF, self).__init__()
    self.conv1 = nn.Conv2d(1, 32, 3, 1)
    self.conv2 = nn.Conv2d(32, 64, 3, 1)
    self.dropout1 = nn.Dropout2d(0.25)
    self.dropout2 = nn.Dropout2d(0.5)
    self.fc1 = nn.Linear(9216, 128)
    self.fc2 = nn.Linear(128, 10)

  def forward(self, x):
    x = self.conv1(x)
    x = F.relu(x)
    x = self.conv2(x)
    x = F.relu(x)
    x = F.max_pool2d(x, 2)
    x = self.dropout1(x)
    x = torch.flatten(x, 1)
    x = self.fc1(x)
    x = F.relu(x)
    x = self.dropout2(x)
    x = self.fc2(x)
    return x

```

```

✓ 0 [16] class NetF1(nn.Module):
  def __init__(self):
    super(NetF1, self).__init__()
    self.conv1 = nn.Conv2d(1, 16, 3, 1)
    self.conv2 = nn.Conv2d(16, 32, 3, 1)
    self.dropout1 = nn.Dropout2d(0.25)
    self.dropout2 = nn.Dropout2d(0.5)
    self.fc1 = nn.Linear(4608, 64)
    self.fc2 = nn.Linear(64, 10)

  def forward(self, x):
    x = self.conv1(x)
    x = F.relu(x)
    x = self.conv2(x)
    x = F.relu(x)
    x = F.max_pool2d(x, 2)
    x = self.dropout1(x)
    x = torch.flatten(x, 1)
    x = self.fc1(x)
    x = F.relu(x)
    x = self.dropout2(x)
    x = self.fc2(x)
    return x

```

2. Переопределяем функцию обучения и тестирования:

```

[17] def fit(model,device,optimizer,scheduler,criterion,train_loader,val_loader,Temp,epochs):
    data_loader = {'train':train_loader,'val':val_loader}
    print("Fitting the model...")
    train_loss,val_loss=[],[]
    for epoch in range(epochs):
        loss_per_epoch,val_loss_per_epoch=0,0
        for phase in ('train','val'):
            for i,data in enumerate(data_loader[phase]):
                input,label = data[0].to(device),data[1].to(device)
                output = model(input)
                output = F.log_softmax(output/Temp,dim=1)
                #calculating loss on the output
                loss = criterion(output,label)
                if phase == 'train':
                    optimizer.zero_grad()
                    #grad calc w.r.t Loss func
                    loss.backward()
                    #update weights
                    optimizer.step()
                    loss_per_epoch+=loss.item()
                else:
                    val_loss_per_epoch+=loss.item()
            scheduler.step(val_loss_per_epoch/len(val_loader))
        print("Epoch: {} Loss: {} Val_Loss: {}".format(epoch+1,loss_per_epoch/len(train_loader),val_loss_per_epoch/len(val_loader)))
        train_loss.append(loss_per_epoch/len(train_loader))
        val_loss.append(val_loss_per_epoch/len(val_loader))
    return train_loss,val_loss

def test(model,device,test_loader,epsilon,Temp,attack):
    correct=0
    adv_examples = []
    for data, target in test_loader:
        data, target = data.to(device), target.to(device)
        data.requires_grad = True
        output = model(data)
        output = F.log_softmax(output/Temp,dim=1)
        init_pred = output.max(1, keepdim=True)[1]
        if init_pred.item() != target.item():
            continue
        loss = F.nll_loss(output, target)
        model.zero_grad()
        loss.backward()
        data_grad = data.grad.data
        if attack == "fgsm":
            perturbed_data = fgsm_attack(data,epsilon,data_grad)
        elif attack == "ifgsm":
            perturbed_data = ifgsm_attack(data,epsilon,data_grad)
        elif attack == "mifgsm":
            perturbed_data = mifgsm_attack(data,epsilon,data_grad)
        output = model(perturbed_data)
        final_pred = output.max(1, keepdim=True)[1]
        if final_pred.item() == target.item():
            correct += 1
            if (epsilon == 0) and (len(adv_examples) < 5):
                adv_ex = perturbed_data.squeeze().detach().cpu().numpy()
                adv_examples.append( (init_pred.item(), final_pred.item(), adv_ex) )
            else:
                if len(adv_examples) < 5:
                    adv_ex = perturbed_data.squeeze().detach().cpu().numpy()
                    adv_examples.append( (init_pred.item(), final_pred.item(), adv_ex) )
        final_acc = correct/float(len(test_loader))
    print("Epsilon: {} \t Test Accuracy = {} / {} = {}".format(epsilon, correct, len(test_loader), final_acc))
    return final_acc,adv_examples

```

3. Создаем функцию защиты методом дистилляции:

```

def defense(device,train_loader,val_loader,test_loader,epochs,Temp,epsilons):
    modelF = NetF().to(device)
    optimizerF = optim.Adam(modelF.parameters(),lr=0.0001, betas=(0.9, 0.999))
    schedulerF = optim.lr_scheduler.ReduceLROnPlateau(optimizerF, mode='min', factor=0.1, patience=3)
    modelF1 = NetF1().to(device)
    optimizerF1 = optim.Adam(modelF1.parameters(),lr=0.0001, betas=(0.9, 0.999))
    schedulerF1 = optim.lr_scheduler.ReduceLROnPlateau(optimizerF1, mode='min', factor=0.1, patience=3)
    criterion = nn.NLLLoss()
    lossF,val_lossF=fit(modelF,device,optimizerF,schedulerF,criterion,train_loader,val_loader,Temp,epochs)
    fig = plt.figure(figsize=(5,5))
    plt.plot(np.arange(1,epochs+1), lossF, "--",label="Loss")
    plt.plot(np.arange(1,epochs+1), val_lossF,"o-",label="Val Loss")
    plt.title("Network F")
    plt.xlabel("Num of epochs")
    plt.legend()
    plt.show()
    #converting target labels to soft labels
    for data in train_loader:
        input, label = data[0].to(device),data[1].to(device)
        softlabel = F.log_softmax(modelF(input),dim=1)
        data[1] = softlabel
    lossF1,val_lossF1=fit(modelF1,device,optimizerF1,schedulerF1,criterion,train_loader,val_loader,Temp,epochs)
    fig = plt.figure(figsize=(5,5))
    plt.plot(np.arange(1,epochs+1), lossF1, "--",label="Loss")
    plt.plot(np.arange(1,epochs+1), val_lossF1,"o-",label="Val Loss")
    plt.title("Network F'")
    plt.xlabel("Num of epochs")
    plt.legend()
    plt.show()
    model = NetF1().to(device)
    model.load_state_dict(modelF1.state_dict())
    for attack in ("fgsm","ifgsm","mifgsm"):
        accuracies = []
        examples = []
        for eps in epsilons:
            acc, ex = test(model,device,test_loader,eps,"fgsm")
            accuracies.append(acc)
            examples.append(ex)
        plt.figure(figsize=(5,5))
        plt.plot(epsilons, accuracies, "--")
        plt.title(attack)
        plt.xlabel("Epsilon")
        plt.ylabel("Accuracy")
        plt.show()
        cnt = 0
        plt.figure(figsize=(8,10))
        for i in range(len(epsilons)):
            for j in range(len(examples[i])):
                cnt += 1
                plt.subplot(len(epsilons),len(examples[0]),cnt)
                plt.xticks([], [])
                plt.yticks([], [])
                if j == 0:
                    plt.ylabel("Eps: {}".format(epsilons[i]), fontsize=14)
                orig,adv,ex = examples[i][j]
                plt.title("{} -> {}".format(orig, adv))
                plt.imshow(ex, cmap="gray")
        plt.tight_layout()
        plt.show()

```

4. Получаем результаты оценки защиты сетей:


```

Temp=100
epochs=10
epsilons=[0,0.007,0.01,0.02,0.03,0.05,0.1,0.2,0.3]
defense(device,train_loader,val_loader,test_loader,epochs,Temp,epsilons)

```



Fitting the model...

Epoch	Loss	Val_Loss
1	0.5481338738800597	1.003228621557355e-05
2	0.3353229173747994	0.00045950331687927244
3	0.26381421273058636	0.00011647366284596501
4	0.2061726351569701	0.00046973809078335765
5	0.16251364960605083	4.2924879031488676e-06
6	0.13502419098830928	1.265520486049354e-06
7	0.11464148852530506	1.9935742660891264e-06
8	0.10202881376955898	8.956958248745651e-06
9	0.09666660607148457	1.391292948028422e-07
10	0.09198299259128082	2.5270764308515935e-08


































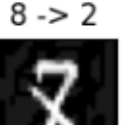









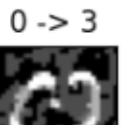


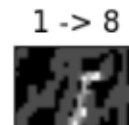

Fitting the model...

Epoch	Loss	Val_Loss
1	0.6784468315404466	7.445180974900722e-05
2	0.4882124064700297	1.9082169979810715e-05
3	0.4424660353774885	1.618420109152794e-05
4	0.40715175550074595	6.557496637105942e-05
5	0.3757558067303607	0.00021618845984339714
6	0.3397686602646741	2.1136321491212584e-05
7	0.2946280665979756	7.914992245787288e-06
8	0.260406978453543	2.3180012707598506e-06
9	0.23156262885010914	2.070512989848794e-06
10	0.20667021672648633	1.3112946021465177e-09



Epsilon	Test Accuracy
0	9119 / 10000 = 0.9119
0.007	9147 / 10000 = 0.9147
0.01	9026 / 10000 = 0.9026
0.02	9036 / 10000 = 0.9036
0.03	8931 / 10000 = 0.8931
0.05	8782 / 10000 = 0.8782
0.1	8024 / 10000 = 0.8024
0.2	4639 / 10000 = 0.4639
0.3	1817 / 10000 = 0.1817
0	9119 / 10000 = 0.9119
0.007	9063 / 10000 = 0.9063
0.01	9114 / 10000 = 0.9114
0.02	8996 / 10000 = 0.8996
0.03	8926 / 10000 = 0.8926
0.05	8778 / 10000 = 0.8778
0.1	7991 / 10000 = 0.7991
0.2	4662 / 10000 = 0.4662
0.3	1852 / 10000 = 0.1852
0	9113 / 10000 = 0.9113
0.007	9069 / 10000 = 0.9069
0.01	9041 / 10000 = 0.9041
0.02	9033 / 10000 = 0.9033
0.03	8945 / 10000 = 0.8945
0.05	8763 / 10000 = 0.8763
0.1	8084 / 10000 = 0.8084
0.2	4628 / 10000 = 0.4628
0.3	1853 / 10000 = 0.1853



Eps: 0	6 -> 6 	3 -> 3 	5 -> 5 	7 -> 7 	8 -> 8 
Eps: 0.007	1 -> 1 	1 -> 1 	9 -> 9 	2 -> 2 	9 -> 9 
Eps: 0.01	8 -> 8 	5 -> 5 	2 -> 2 	8 -> 8 	4 -> 4 
Eps: 0.02	8 -> 8 	7 -> 7 	0 -> 0 	2 -> 2 	1 -> 1 
Eps: 0.03	2 -> 2 	5 -> 5 	9 -> 9 	9 -> 9 	2 -> 2 
Eps: 0.05	6 -> 6 	4 -> 9 	0 -> 0 	9 -> 9 	9 -> 9 
Eps: 0.1	8 -> 2 	6 -> 6 	6 -> 6 	0 -> 5 	2 -> 2 
Eps: 0.2	8 -> 7 	7 -> 5 	6 -> 6 	2 -> 0 	0 -> 0 
Eps: 0.3	0 -> 3 	1 -> 8 	6 -> 7 	1 -> 8 	1 -> 3 

Заключение

В ходе выполнения лабораторной и практической работы были успешно реализованы атаки на защищенные и незащищенные модели НС, показавшие разную эффективность. Способ дистилляционной защиты смог предоставить достаточно высокий уровень безопасности в сравнении с незащищенными моделями. Так, точность оказалась на уровне 91%.